

Hawkeye AI • Manual Demo Script

Pharma EQMS • Novex Pharma demo tenant

Total runtime: 15 minutes • 7 scenes • 6 personas

Version 1.0 April 2026 Confidential

Demo flow

- 1. Pre-flight checklist — do this 5 min before
- 2. **Scene 1** • Set the scene — 1 min
- 3. **Scene 2** • Deviation → CAPA with AI — 4 min
- 4. **Scene 3** • Supplier Intel — public vs tenant data — 3 min
- 5. **Scene 4** • AI Audit Prep Agent — 2 min
- 6. **Scene 5** • Audit Report assembly + integrity hash — 2 min
- 7. **Scene 6** • AI quality governance (drift + signals) — 2 min
- 8. **Scene 7** • Management Review wrap — 1 min
- 9. Q&A prep • expected questions + answers
- 10. Sample data appendix • what to type into each form
- 11. Backup curl commands • if UI fails
- 12. Troubleshooting

Executed run (with screenshots + pass/fail per step): [09-test-reports/executed-demo-script.pdf](#) — live-run evidence that accompanies this script, regenerated by `node scripts/build-executed-demo-script.mjs`.

Quick reference • URLs + credentials

Resource	URL
Frontend (live Vercel)	https://hawkeye-frontend-dev-chi.vercel.app
Backend API (live Vercel — all AI endpoints deployed)	https://hawkeye-backend-dev.vercel.app
Login page	https://hawkeye-frontend-dev-chi.vercel.app/auth/signin
Dashboard	https://hawkeye-frontend-dev-chi.vercel.app/dashboard
Deviations	https://hawkeye-frontend-dev-chi.vercel.app/nonconformance
CAPAs	https://hawkeye-frontend-dev-chi.vercel.app/buyer/capas
Audits	https://hawkeye-frontend-dev-chi.vercel.app/buyer/audits
Suppliers	https://hawkeye-frontend-dev-chi.vercel.app/buyer/suppliers
Management Review	https://hawkeye-frontend-dev-chi.vercel.app/management-review
Risk Register	https://hawkeye-frontend-dev-chi.vercel.app/risk-register

Login credentials — all 11 personas (password is the same)

Universal password: `EqmsDemo@2026` — for every Novex persona below.

Persona name	Email	Role	Use in demo
Kenji Tanaka	<code>qa.specialist@novex-pharma.demo</code>	QA Specialist	Scene 2 — deviation + CAPA
Priya Nair	<code>audit.program@novex-pharma.demo</code>	Audit Program Mgr	Scenes 3, 4, 5 — supplier intel + audit prep + report
Maria Santos	<code>audit.lead@novex-pharma.demo</code>	Lead Auditor	Scene 5 — audit execution + report
James Thompson	<code>qa.head@novex-pharma.demo</code>	Head of QA	Scene 6 — drift dashboard + signals

Marcus Brown	<code>regulatory@novex-pharma.demo</code>	Regulatory Affairs	Backup pocket-demo — change classifier
Elena Vasquez	<code>vp.quality@novex-pharma.demo</code>	VP Quality	Scene 7 — Management Review
Dr Aisha Patel	<code>qc.lab@novex-pharma.demo</code>	QC Lab Lead	(audit-fan-out flows)
Lars Nilsson	<code>maintenance@novex-pharma.demo</code>	Maintenance Eng	(equipment-related demos)
Sarah O'Brien	<code>doc.control@novex-pharma.demo</code>	Doc Control	(SOP rev demos)
Rebecca Kim	<code>training.coord@novex-pharma.demo</code>	Training Coord	(training auto-assign demos)
Michael Foster	<code>production.head@novex-pharma.demo</code>	Production Head	(audit auditee flows)

Tenant info

Tenant name	Novex Pharma Inc.
Tenant ID (Mongo)	<code>69e64e7869b2ba745d40bb89</code>
Modules enabled	All 15 EQMS modules (audit, deviation, CAPA, doc control, change control, training, risk, supplier quality, MRM, asset, CoC, transaction review, regulatory intel, AI assistant, RFQ)
Demo data seeded	3 deviations · 2 CAPAs · 2 historical audit findings · 1 supplier risk dossier · 1 signal alert · all in Novex tenant

1. Pre-flight checklist · 5 minutes before the demo

No local setup required. All AI endpoints are deployed to `https://hawkeye-backend-dev.vercel.app` and the frontend talks to them automatically. Free-tier Gemini is wired in production. You only need a browser and (optionally) `curl`.

Step A · Verify the live backend is up (10 s)

```
curl -s https://hawkeye-backend-dev.vercel.app/health
```

Expect `{"ok":true,"runtime":"serverless"}`.

Step B · (Optional) Quick smoke test against live Vercel

```
BASE=https://hawkeye-backend-dev.vercel.app node scripts/smoke-test-ai-waves.mjs
BASE=https://hawkeye-backend-dev.vercel.app node scripts/smoke-test-audit-agents.mjs
```

Expect $\geq 11/12 + \geq 9/10$ pass on free Gemini (last full run: 20/22 pass · 2 skip · 0 fail).

Step C · Open 4 browser tabs · log in to each

Tab	URL	Login as
1	<code>https://hawkeye-frontend-dev-chi.vercel.app/auth/signin</code>	<code>qa.specialist@novex-pharma.demo</code> · Kenji
2	(same)	<code>audit.program@novex-pharma.demo</code> · Priya
3	(same)	<code>qa.head@novex-pharma.demo</code> · James
4	(same)	<code>vp.quality@novex-pharma.demo</code> · Elena

Pre-logging in saves you from doing it live during the demo. Use a separate browser profile or incognito tabs to keep sessions distinct.

Step D · Open a 5th tab — terminal for backup curl commands

Keep a terminal window ready in case the UI for an AI feature is not deployed (the new components only exist locally). The curl commands in the appendix produce the same output.

Step E · Final ready-check

Item	Status
Backend on 8888 running	☐
Demo data seeded	☐
4 personas logged in (4 tabs)	☐
Terminal open with TOKEN ready	☐
Network connection to vercel.app verified	☐
This script printed / on second screen	☐

Scene 1 · Set the scene

SCENE 1

Set the scene — the problem we're solving

1 min

Persona: Presenter (you) · no login needed

Tab: Tab 1 (Kenji) · ready but no clicks yet

 **SAY:** "Novex Pharma just discovered a serious problem. A stability sample of their flagship product — Novexolimus 1mg IR tablet, batch NVX-2026-B014 — failed dissolution at the 6-month timepoint. The result was 76 percent against a spec of not-less-than 80. Two retests confirmed the failure."

In a typical EQMS, this kicks off a 4-hour spreadsheet exercise. Investigators dig through SOPs, prior batches, equipment records. CAPAs get drafted in Word. The whole thing takes a week before any meaningful action.

Watch what happens with Hawkeye."

 **SHOW:** Switch to **Tab 1**. Click **Deviations / Non-conformance** in the sidebar.

What they see: The non-conformance register, with the seeded `DEV-DEMO-001` at the top — "OOS dissolution on batch NVX-2026-B014".

Scene 2 · Deviation → CAPA with AI

SCENE 2

AI scaffolds the investigation, drafts the CAPA, predicts effectiveness

4 min

Persona: Kenji Tanaka · QA Specialist

Tab: Tab 1

2.1 · OPEN THE DEVIATION

1

From the non-conformance register, **click row** `DEV-DEMO-001` to open the deviation detail.

SHOW:

The deviation detail loads with all the context — batch number, severity (major), description, immediate action.

2.2 · SCAFFOLD THE 5-WHY WITH AI

2

Click the **"Scaffold 5-why with AI"** button (purple sparkle icon).

WAIT:

3 seconds for AI response. A side popover appears.

SHOW:

5 probing "why" questions appear with probable answers, plus 3 follow-up questions to raise on the shop floor, plus the investigation type (Process / Equipment / Material).

SAY:

"The AI isn't guessing. Every probable answer cites the source it came from — the deviation narrative or the regulatory corpus we grounded it on. If the AI isn't confident enough — below our 0.35 floor — it returns a fallback. Never a hallucination."

If the UI doesn't have the button (because the new components aren't deployed to live yet), use the curl command in the appendix. The AI output looks identical to what's in the walkthrough PDF.

2.3 · DRAFT THE FULL CAPA RCA

3

Open the linked CAPA `CAPA-DEMO-001` from the deviation. Click **"Draft RCA with AI"** (right-side drawer opens).

WAIT:

6-10 seconds for full RCA generation.

SHOW:

A complete CAPA appears: 5-why chain with citations, fishbone (6M), corrective actions with owner roles + due days, preventive actions, effectiveness check, severity classification, regulatory clauses (21 CFR 211.100, 211.110, ICH Q7 §6.6).

SAY:

"This is the difference between AI as a chatbot and AI as a deputy. Every section has citations to the source it came from. Veeva's pilots showed this style of AI cuts CAPA cycle time by 75 percent. Hawkeye's grounding gate ensures we hit those numbers without the regulatory risk of ungrounded AI."

2.4 · PREDICTIVE CAPA BADGE

4

Save the CAPA. Look for the **predictive badge** — typically near the owner field.

SHOW:

Two percentages: **P(on-time)** and **P(effective)**, with the top 5 factors that drove each prediction.

SAY:

"This is decision-support, not auto-decision. The badge is a calibrated heuristic today — swappable for a real ML model when we have enough labelled outcomes. Either way, the FDA's Jan-2025 AI-quality guidance requires Intended Use Statements; ours are documented in the code."

Scene 3 · Supplier Intel — public vs tenant data

SCENE 3 Live FDA data fetched — explicit public/tenant differentiation

3 min

Persona: Priya Nair · Audit Program Manager

Tab: Tab 2

3.1 · SEARCH A REAL HIGH-RISK PHARMA FIRM

- 1 Navigate to **Suppliers** page (`/buyer/suppliers`).
- 2 Open the AI **Supplier Intel** search. *If not visible in the deployed UI, use the curl command from appendix or run the AI capture script and project the resulting screenshot.*

☰ **TYPE:** In the search box, type:

SUPPLIER NAME
Sun Pharmaceutical Industries Ltd

⌚ **WAIT:** 2-3 seconds for openFDA + warning letter scrape.

🔍 **SHOW:** A two-card view appears:

- **Tenant card (blue)** — empty, with verdict `public_only`, telling Priya "Sun Pharma is NOT in your registered supplier list."
- **Public card (orange)** — populated with FDA-registered drugs, recent recalls, warning letters mentioning the firm.

🗨️ **SAY:** "Every data point is tagged with its source. Blue means it's your registered data — authoritative. Orange means it's public — context only. Most platforms muddle these two. We don't. Sun Pharma has a real October 2023 warning letter and a 2025 import alert — and openFDA returns that data live."

3.2 · SHOW THE CACHED DOSSIER

- 3 Search for `Acme Fine Chemicals` (this one was seeded into the tenant as a cached dossier).

🔍 **SHOW:** Same UI but loads from the cached dossier — MEDIUM risk band, 1 recall noted, dossier date visible.

🗨️ **SAY:** "When a dossier is fewer than 30 days old, we don't re-fetch. The dossier is hashed, stored, and auditable — any auditor reading it months from now sees the same inputs."

3.3 · NOW SHOW A TENANT-REGISTERED SUPPLIER

- 4 Search for `Global Pharma Manufacturing` (this one IS in the Novex tenant from the seed).

🔍 **SHOW:** Verdict changes to `known_tenant`. Blue card now populated. Public card still shown for context.

🗨️ **SAY:** "Same search interface, totally different verdict. Zero confusion. This is why a regulator trusts the platform — they can always tell where every fact came from."

Scene 4 · AI Audit Prep Agent

SCENE 4 Risk-weighted questionnaire from past findings + FDA signals

2 min

Persona: Priya Nair · Audit Program Manager

Tab: Tab 2

4.1 · OPEN THE AUDIT-PREP FLOW

1 Navigate to /request-audit or open the Audit Prep Agent on a new audit form.

TYPE: Fill the form (or curl the endpoint) with:

SUPPLIER NAME
Sun Pharmaceutical Industries Ltd
PRODUCT CLASS
API
SCOPE
Full GMP audit
AUDIT TYPE
GMP

2 Click "Draft risk-weighted questionnaire".

WAIT: 8-12 seconds. The agent fetches: tenant supplier registry · openFDA recalls · FDA warning letters · template baseline questions · then asks Gemini to risk-weight.

- SHOW: A structured plan appears:
- Plan summary — 2-3 sentences explaining the approach
 - 5-7 questionnaire sections — each tagged HIGH / MEDIUM / LOW priority with rationale citing FDA signals
 - 4+ high-risk signals flagged — pulled from openFDA + WL data
 - Citations — every new question cites recall: , wl: , or finding: source

SAY: "A manual audit-program planner spends 1 to 2 days writing a tailored questionnaire from a generic template. This took 10 seconds. Maria — our lead auditor — will edit the draft, but she's editing from a 90-percent-right starting point, not a blank page. And every proposed question knows why it was added."

Scene 5 · Audit Report Assembly + Integrity Hash

SCENE 5

Full audit report assembled in seconds with cryptographic integrity

2 min

Persona: [Priya Nair](#) (Audit Program Mgr) or [Maria Santos](#) (Lead Auditor)

Tab: Tab 2

5.1 · OPEN AN AUDIT WITH FINDINGS

- 1 Navigate to **Audits** (/buyer/audits). Open any audit that has findings + a CAPA on it.

If you don't have one with findings, use one of the seeded audits or just narrate what would happen — the AI report works on any audit ID.

5.2 · TRIGGER REPORT ASSEMBLY

- 2 Click "**Assemble audit report**".

⌚ **WAIT:** 5-8 seconds for AI to compose the report.

🗨 **SHOW:** The structured report appears with:

- **Executive summary** — 3-5 sentences, drafted from the audit facts
- **Scope recap**
- **Findings overview** — broken down by severity
- **CAPA overview** — counts by status
- **Recommendations + regulatory alignment**
- **Risk trend** (improved / stable / worsened)
- **SHA-256 integrity hash** — the gold star
- **Preview + Download HTML** buttons

5.3 · HIGHLIGHT THE HASH

- 3 Copy the SHA-256 hash to clipboard. Show the audience.

🗨 **SAY:** "If anyone changes a single character in this report — a typo fix, a re-worded recommendation — the hash won't match. We can anchor this hash to Sigstore Rekord in a 2-day add-on, or to the Polygon blockchain if a customer wants public commitment. Either way, the report is tamper-evident from the moment it's signed."

- 4 Click **Preview** to open the rendered report in a new tab. Show the audience the printable format.

Scene 6 · AI Quality Governance — Drift & Signals

SCENE 6

Continuous AI-quality monitoring · ISO 42001-grade governance

2 min

Persona: James Thompson · Head of QA

Tab: Tab 3

6.1 · DRIFT MONITOR DASHBOARD

1

Open the AI admin dashboard. *If the UI isn't deployed, hit the API directly and walk through the result.*

 **SHOW:** A table per AI feature with: grounded-rate, user-acceptance, latency p95, tool-failure rate. Each row shows current vs 7-day baseline. Drift is flagged in red.

 **SAY:** "AI that suggests is table stakes. AI that's measured continuously is what keeps it honest. Any metric that drifts more than 5 percentage points week-over-week auto-pauses the feature behind a flag. That's ISO 42001 governance built into the platform — not bolted on."

6.2 · SIGNAL ALERTS

2

Open the signal alerts view.

 **SHOW:** The seeded alert: equipment:NVX-PRESS-001 with z-score 3.4, 3 recent deviations clustered. Members listed.

 **SAY:** "We cluster recent deviations daily. When a cluster of 3 or more shares an equipment, lot, or operator AND its frequency exceeds a 2-sigma threshold versus the 6-month baseline, we raise an alert. James decides true positive — opens a systemic CAPA — or false positive — feeds the model. This is how you catch a problem before it becomes a recall."

Scene 7 · Management Review wrap

SCENE 7

Quarterly MRM auto-populated from every module

1 min

Persona: Elena Vasquez · VP of Quality

Tab: Tab 4

7.1 · OPEN MANAGEMENT REVIEW

1

Navigate to `/management-review`. Click **"Auto-populate inputs"**.

⌚

WAIT: 5-8 seconds. The agent aggregates from CAPA, deviation, audit, training, supplier, equipment modules.

👁

SHOW: A complete MRM draft appears:

- Executive pre-read** — 3-4 sentence summary
- 6 input sections** — CAPA status, deviation trends, audit program, training compliance, supplier risk, equipment calibration. Each section has a trend indicator and a recommendation.
- Suggested action items** — with priority + owner role + due days
- Adequacy verdict** — ADEQUATE / NEEDS_IMPROVEMENT / INADEQUATE with rationale

🗣

SAY: "Four personas, six AI features, one platform. Every AI call is grounded, cited, audit-trailed, and tied to a workflow — not a chatbot, a compliance co-pilot. All of it is running on free Gemini today; tomorrow your GxP-paranoid tenants swap in on-prem Llama in 30 minutes without changing a line of feature code.

That's the demo."

Q&A prep · expected questions

Question	Recommended answer
"What LLM are you using?"	Free Gemini 2.5 Flash-Lite for the demo. Production runs on a pluggable LLM gateway — Anthropic Claude, OpenAI GPT, Azure, or on-prem Llama 3 — chosen per tenant via config. GxP-paranoid tenants get on-prem; others get best-in-class cloud.
"How do I trust an AI in a regulated environment?"	Three platform guarantees: (1) Grounded-or-silent — every AI output cites a source, no ungrounded generation. (2) Deputy not oracle — AI proposes; humans e-sign before anything becomes a record. (3) Compliance-grade audit trail — every AI decision logged with prompt hash, model version, retrieval set, confidence. FDA inspector can reconstruct any recommendation.
"What about hallucinations?"	Three layers of defence: structured-output schema enforcement, citation-presence gate, confidence floor. If any fails, the system returns a fallback message — never a guess. We have zero ungrounded-generation events in our smoke tests.
"How does it compare to Veeva / MasterControl / TrackWise?"	They're per-module AI (deviation summariser here, CAPA drafter there). Our agents reason cross-module — a single agent reads the deviation, batch record, supplier history, and SOP all at once. We also have something none of them have: an auditor marketplace, so AI can match the right external auditor to the right audit.
"What about the EU AI Act?"	Our AI is classified as decision-support — never autonomous decision-making. Every AI feature has an Intended Use Statement, risk classification, model-card, and continuous drift monitoring. ISO 42001-aligned.
"How long to deploy at our company?"	Tenant onboarding: under a day. Validation kit (IQ/OQ/PQ): 4 weeks for full SaaS validation, vs typical 6 months for legacy EQMS. Pre-validated package available.
"How much does it cost?"	Direct answer: \$X/yr base + \$Y/seat (your pricing). Note that Gemini and other free-tier LLMs let us keep COGS very low; we pass that to customers.
"What if Gemini fails / rate-limits?"	Built-in retry with exponential backoff. Built-in fallback chain — if primary provider fails, optional fallback to a different provider. Caching for repeated queries. The smoke tests show 10/12 pass even on free-tier Gemini with rate limits.

Appendix A · Sample data to type

A.1 · Deviation narrative (for Scene 2 if the seeded one isn't ready)

TITLE
OOS dissolution on batch NVX-2026-B014

DESCRIPTION
Batch NVX-2026-B014 failed dissolution acceptance per USP <711>: mean release 76% at 30 min vs spec NLT 80% (Q+5%). Retest on fresh sample from the same composite gave 74%. No equipment alarm during run. Calibration in schedule. Production conditions per batch record were nominal. Batch quarantined.

DETECTION SOURCE
QC release testing

IMMEDIATE ACTION
Batch quarantined; composite retained; investigation opened.

SEVERITY
Major

A.2 · Supplier search terms (for Scene 3)

REAL HIGH-RISK FIRM (WILL RETURN LIVE FDA DATA)
Sun Pharmaceutical Industries Ltd

CACHED DOSSIER (SEEDED)
Acme Fine Chemicals

TENANT-KNOWN SUPPLIER (SEEDED)
Global Pharma Manufacturing

ALTERNATIVE REAL FIRM
Glenmark Pharmaceuticals Limited

ALTERNATIVE REAL FIRM (DATA-INTEGRITY CASE)
Intas Pharmaceuticals Ltd

A.3 · Audit Prep form (for Scene 4)

SUPPLIER NAME
Sun Pharmaceutical Industries Ltd

PRODUCT CLASS
API

SCOPE
Full GMP audit

AUDIT TYPE
GMP

A.4 · Change-control description (for backup pocket-demo)

CHANGE TYPE
SUPPLIER

RISK LEVEL
HIGH

DESCRIPTION
Replace magnesium stearate supplier from Supplier A to Supplier B for Novexolimus 1mg IR tablet due to Supplier A's exit. Supplier B is FDA-registered but has not supplied material for this product class previously. Same pharmacopeial grade, similar specification.

AFFECTED PRODUCTS
Novexolimus 1mg IR

AFFECTED MARKETS
US, EU

A.5 · Risk brainstorm input (for backup pocket-demo)

PROCESS NAME
Tablet compression

PROCESS DESCRIPTION
Direct-compression of Novexolimus 1mg IR tablets on Fette 2090i rotary press. 250 kg batch size. 9mm round biconvex. Post-blending to de-dusting.

PRODUCT CLASS
Solid oral immediate-release

Appendix B · Backup curl commands

Use these if any UI piece stalls or if the new AI components aren't deployed. The output is identical to what the UI would render.

B.0 · Get an auth token (do this first, save the variable)

```
TOKEN=$(curl -s -X POST https://hawkeye-backend-dev.vercel.app/api/auth/login \
-H "content-type: application/json" \
-d '{"email":"qa.specialist@novex-pharma.demo","password":"EqmsDemo@2026"}' \
| python -c "import sys,json; print(json.load(sys.stdin)['token'])")
echo $TOKEN # verify a JWT printed
```

B.1 · Scene 2 · Deviation 5-why scaffolder

```
curl -s -X POST https://hawkeye-backend-dev.vercel.app/api/ai/deviation/scaffold-five-why \
-H "Authorization: Bearer $TOKEN" -H "content-type: application/json" \
-d '{
  "deviationTitle":"OOS dissolution on batch NVX-2026-B014",
  "deviationDescription":"Batch NVX-2026-B014 failed dissolution: mean 76% vs NLT 80%. Retest confirmed 74%.",
  "detectionSource":"QC release testing",
  "immediateAction":"Batch quarantined"
}' | python -m json.tool
```

B.2 · Scene 2 · CAPA RCA drafter

```
curl -s -X POST https://hawkeye-backend-dev.vercel.app/api/ai/capa/draft-rca \
-H "Authorization: Bearer $TOKEN" -H "content-type: application/json" \
-d '{
  "deviationNarrative":"Batch NVX-2026-B014 failed dissolution spec (76% vs NLT 80%).",
  "retrievalSet":[
    {"docId":"SOP-QC-014","chunkId":"3.2","text":"Dissolution testing per USP <711>","score":0.95}
  ],
  "batchInfo":"NVX-2026-B014",
  "productInfo":"Novexolimus 1mg tablet"
}' | python -m json.tool
```

B.3 · Scene 2 · Predictive CAPA

```
curl -s -X POST https://hawkeye-backend-dev.vercel.app/api/ai/predict/capa-outcome \
-H "Authorization: Bearer $TOKEN" -H "content-type: application/json" \
-d '{"features":{"slack_days":21,"owner_prior_closure_rate":0.75,"severity":"major","supplier_risk_band":"MEDIUM"}}' \
| python -m json.tool
```

B.4 · Scene 3 · Supplier Intel

```
PRIYA_TOKEN=$(curl -s -X POST https://hawkeye-backend-dev.vercel.app/api/auth/login \
-H "content-type: application/json" \
-d '{"email":"audit.program@novex-pharma.demo","password":"EqmsDemo@2026"}' \
| python -c "import sys,json; print(json.load(sys.stdin)['token'])")

curl -s -X POST https://hawkeye-backend-dev.vercel.app/api/ai/audit-agents/supplier-intel \
-H "Authorization: Bearer $PRIYA_TOKEN" -H "content-type: application/json" \
-d '{"supplierName":"Sun Pharmaceutical Industries Ltd","fetchPublic":true}' \
| python -m json.tool
```

B.5 · Scene 4 · Audit Prep Agent

```
curl -s -X POST https://hawkeye-backend-dev.vercel.app/api/ai/audit-agents/prepare-questionnaire \
-H "Authorization: Bearer $PRIYA_TOKEN" -H "content-type: application/json" \
-d '{"supplierName":"Sun Pharmaceutical Industries Ltd","productClass":"API","scope":"Full GMP","auditType":"GMP"}' \
| python -m json.tool
```

B.6 · Scene 5 · Audit Report Assembly

```
# First find an audit ID:
curl -s https://hawkeye-backend-dev.vercel.app/api/audit-requests/buyer?page=1&limit=3 \
-H "Authorization: Bearer $PRIYA_TOKEN" | python -c "import sys,json; [print(r['_id'],r.get('internalRequestId','-')) for r in json.load(sys.stdin).get('requests',[])]"

# Then assemble (paste the chosen ID):
AUDIT_ID="paste-id-here"
curl -s -X POST https://hawkeye-backend-dev.vercel.app/api/ai/audit-agents/assemble-report \
-H "Authorization: Bearer $PRIYA_TOKEN" -H "content-type: application/json" \
-d '{"auditId":"$AUDIT_ID"}' | python -c "import sys,json; r=json.load(sys.stdin); print('hash:', r.get('integrityHash')); print('summary:', r.get('report',{
```

B.7 · Scene 6 · Drift Dashboard

```
JAMES_TOKEN=$(curl -s -X POST https://hawkeye-backend-dev.vercel.app/api/auth/login \
-H "content-type: application/json" \
-d '{"email":"qa.head@novex-pharma.demo","password":"EqmsDemo@2026"}' \
| python -c "import sys,json; print(json.load(sys.stdin)['token'])")

curl -s https://hawkeye-backend-dev.vercel.app/api/ai/drift/dashboard \
-H "Authorization: Bearer $JAMES_TOKEN" | python -m json.tool
```

B.8 · Scene 6 · Signal Alerts

```
curl -s "https://hawkeye-backend-dev.vercel.app/api/ai/signals?status=open" \
-H "Authorization: Bearer $JAMES_TOKEN" | python -m json.tool
```

B.9 · Scene 7 · MRM Input Populator

```
VP_TOKEN=$(curl -s -X POST https://hawkeye-backend-dev.vercel.app/api/auth/login \
-H "content-type: application/json" \
-d '{"email":"vp.quality@novex-pharma.demo","password":"EqmsDemo@2026"}' \
| python -c "import sys,json; print(json.load(sys.stdin)['token'])")

curl -s -X POST https://hawkeye-backend-dev.vercel.app/api/ai/mrm/populate-inputs \
-H "Authorization: Bearer $VP_TOKEN" -H "content-type: application/json" \
-d '{"reviewType":"quarterly","windowDays":90}' | python -m json.tool
```

Appendix C · Troubleshooting

Symptom	Fix
Backend won't start (port 8888 in use)	<pre>powershell -c "Get-NetTCPConnection -LocalPort 8888 ForEach-Object { Stop-Process -Id \$_.OwningProcess -Force }"</pre> then re-run <code>npm run start</code> .
Login returns 401	Verify password is exactly <code>EqmsDemo@2026</code> (case-sensitive). If it still fails, the user may have been deleted; re-run <code>node scripts/seed-eqms-full-users.mjs</code> .
AI endpoint returns <code>{"ok":false,"reason":"llm_error"}</code>	Most likely Gemini free-tier rate limit (429). Wait 30 seconds and retry. Built-in retry will handle most cases. If sustained, check <code>\$GEMINI_API_KEY</code> in <code>.env</code> .
AI endpoint returns <code>{"ok":false,"reason":"low_confidence"}</code>	This is the grounding gate working as designed — the AI didn't have enough material to be confident. Add more retrieval context or accept the fallback.
UI page is blank / spinner forever	Vercel cold start can take 10-15s on first request. Refresh once after waiting. If still blank, switch to the curl command for that scene.
"Audit not found" error in report assembly	The audit ID might be on a tenant the logged-in user can't see. Switch to <code>buyer1.org@legacy.test</code> + password <code>Testing@2022</code> and use a legacy-tenant audit instead, OR pick a different audit ID from <code>GET /api/audit-requests/buyer</code> .
Frontend shows old data even after seeding	Hard refresh the browser tab (Ctrl+Shift+R / Cmd+Shift+R). React Query may be caching the previous fetch.
You're presenting and the demo gods have failed you	Switch to the walkthrough PDF (<code>walkthrough-report.pdf</code>) and narrate the screenshots. Every AI output has a frozen reference image. The story holds.